



DSC Capstone Sequence

Lecture 08
Long Running Jobs



Lecture Outline

- Remote Job Submission
 - Running Long Running Jobs)
- Test Data
 - Developing Long Running Jobs
- Logging
 - Debugging Long Running Jobs


Remote Job Submission

Long Running Jobs

Why watch your 3HR job run, when you can:

- Log onto a server
- Start your job (build script)
- Log off server
- [...]
- Return when job is finished

Compute time != Keyboard time



Requirements for Long Running Jobs

Effective use of running "hands-off" project execution requires:

- The server to remain running after logout
- Non-interactive code that saves output to file
 - e.g. a build script with targets
- Writing code that's largely correct
 - It is a *long* running job, afterall!
- Good logging
 - Debugging is more difficult and more removed.

DSMLP and Background Pods

- DSMLP is a "Kubernetes Cluster" that manages compute resources (pods).
- User requests a pod (e.g. using a launch script) and Kubernetes allocates resources for you (a pod running a container with requested memory).
- Kubernetes has a command-line tools for interacting with Pods:
 - [kubectrl](#) is the kubernetes "control" tool -- allowing you to list/delete your running pods.
 - kubesh is the kubernetes "shell" -- connects you to the shell in a specified pod.

Kubernetes command-line tools allow you to connect, disconnect, and reconnect to pods without shutting them down!

Connecting to DSMLP Background Pods

- To launch a background pod:
 - `launch-180.sh -G <group> -b`
- Text displayed to the terminal lists your `pod-ID` with instructions.
- Connect to your pod using the command:
 - `kubesh <pod-ID>`
- After starting job, disconnect from pod (`ctrl-d`)
- Check the pod is still running:
 - `kubectl get pods`
- **Delete your pod when finished**
 - `kubectl delete pods`

Running Jobs on Background Pods

- You can now attach/detach to/from a running pod, without shutdown.
- From within the pod, how do you run a job "in the background" so that you can exit the pod with the job still running?

Running a job in the background:

- `python run.py target &`
- ``&`` tells shell to start the process in the background
 - Terminal remains free

To stop a job, use ``kill <PID>`` where PID is the Process ID.

- Find your job's PID using the ``jobs`` command, or ``ps`` command.

Advice for Long Running Jobs

- Be sure code is *correct* before running (else you waste *a lot* of time!)
- Develop code on small data w/small resources before committing to running the big job
- Use a lot of logging; they are difficult to debug.
- Increase the timeout on a POD by editing the launch script:
 - Pods shutdown automatically after ~3 hours
 - Edit the launch script as advised [here](#)
 - Change the K8S_TIMEOUT_SECONDS variable (up to 12 hours; longer with permission).

Test Data

Developing on Test Data

- Create small "realistic" test data on which to develop code
 - Speeds up code development iterations
 - Using real (large) data naturally takes longer...
- Integrate test data into doctests and/or unittests
 - Test changes to your code against this data to make debugging easier.
- Test data works well to check that the steps of your project "fit together" properly.

What is Test Data?

Small made-up data, that is realistic enough to *test* code (quickly)

- It is NOT "real" data; the developer creates it.
 - Test data should NOT be "a small sample" of real data
 - Test data is designed to test correctness of the code.
 - Developer should design each line of test data to test a specific attribute!
- Should not contain sensitive information.
- A test-data file should only be a few lines.
 - You should be able to test the correctness of code on this data BY HAND.
- The statistics of the test data may not look realistic, but the schema should.

Test Data and Versioning (Git)

- Test data *should* be versioned!
- It is developed, and changes, to test code.
 - This meets the criteria for versioning: changing material that is not built from a static data source and versioned code.
- Place test data in a testing directory, like `test/testdata`.
 - `test` is where developers place their unit-tests.

Test data is not real data; **never version control data.**

How to use test data

- Test data tests the correctness of code *before running it on real data*.
 - Run the test data on a small container (small RAM; fast), before starting a large container and letting it run overnight on the "real" data.
- Particularly useful for testing the steps of your pipeline fit together.
- Substitute test-data for real data in your build script to quickly run tests.

test Target

- Implement a `test-data` target that runs the same code as `data`, only on your test data.
- Then, to train a model on your test data, you might run:
 - `python run.py test-data features model`
- The target `test` is a standard target that implements the `all` target on test data. That is, it builds your entire project on test data.

How Project Code is Graded

Your project will be run on test-data on the DSMLP server, in a clean location:

1. `launch-180.sh -i <student-dockerhub-repo>`
2. `git clone <student-repo>` (or copy code from gradescope)
3. `python run.py test`
4. (run script that checks generated files)

In Step 3, the test target should generate all your project output, run on test data that's versioned in your repository. It should take no more than a few minutes.

Logging

Logging

When running build scripts without interaction:

- Hard to tell why something failed
- Takes a long time to try and retry
- You have no access to terminal output
 - Must log messages to file, for later reading

Print Debugging

- Simple: log processes with print statements
- You can "redirect" standard output (stdout) using `>`
 - `python run.py all > log.txt &`
 - This redirects all terminal output to the file `log.txt`
- Disadvantage: hard to tell which statement came from where and when.

logging Module

- logging module comes in the python standard library
- Automatically records:
 - when logging happened (can use this to time your steps!)
 - which file/function the logging occurred in
- Can set "logging levels" to restrict output based on context
 - debug/info/warn/critical
 - use debug only when debugging
 - use info to record performance (e.g. time for each step)

Using the logging module

```
logging.basicConfig(filename=logname, filemode='a',  
                    format='%(asctime)s,%(msecs)d %(name)s %(levelname)s %(message)s',  
                    datefmt='%H:%M:%S',  
                    level=logging.DEBUG)  
  
logging.info("Log info Here")
```